



(a) The envelope, established: $n = 1$




retraining: rank $\times 2.69$, channel $+0.035$
re-export deterministic to 5 s.f. (0.58612 against 0.5861 recorded)


(b) The seven between-arm differences, against it



(c) The asymmetry, beside it



(d) The controlled repeat that would make (a) an interval



F1. The envelope of panel (a) is ONE retraining pair on ONE configuration on ONE stack (D2 arm H, seed 42, identical arguments). It cannot separate rank-specific variance from stack non-determinism, architecture or schedule, and it is drawn as a single observation, not as an estimated distribution: no confidence band is shown because none has been measured. Statistic R1 (Roy & Vetterli order-1, centred), block confound-residualised for every D2 and D1-B point and RAW for the Phase 1b point - and the Phase 1b arms were never verified matched (draft 3.4). Panel (b)'s x-axis is a ratio of two R1 values within one comparison; no two rank statistics share an axis anywhere in this figure. The claim does not rest on this figure alone: F2 reaches the same conclusion from 8 within-arm degrees of freedom without using the 2.69x number. Panels (b) and (c) sit side by side because the channel is quoted as a paired within-run difference and rank is not - there is no paired form of "this run has higher rank".