




(a) The decomposition: how much of each quantity is the ARM, and how much is the training SEED



(b) The twelve artifacts, per arm: rank





... and information, same rows



(c) The single clearest object, enlarged: D2 arm H, seed 44 against seed 43

one arm, no arm contrast to argue about, inside RankMe's reserved scope by RankMe's own sentence



F2. Four arms and three seeds is 8 within-arm degrees of freedom, not a large design. $F(3, 8) = 1.41$ for canonical effective rank is a FAILURE TO REJECT: the arm effect on rank is not resolvable here, which is not the same as there being none. The seed changes nothing about objective, architecture, data, split or schedule - that is what makes it the nuisance factor. Statistic R1 (Roy & Vetterli order-1, centred), block confound-residualised; the RankMe row is uncentred and RAW, per its own definition, and is included because it is the answer to "you evaluated a centred variant" - it is WORSE than ours, not better. Rank-type quantities are decomposed on the log scale and the CCA on the raw scale, as the axis label states; no two rank statistics share an axis. Panel (c) is inside RankMe's own reserved scope: three seeds of one arm are "different runs of a given method". Values recomputed here from the per-artifact metrics and asserted equal to the printed decomposition in data/ws_p2/out/p2_run.log.