


F3 The reproducibility floor is a property of the ARM, not of the statistic same configuration, same global step, only the training seed differs



arm	seed 42	seed 43	seed 44	seed 45	seed 46	spread
programme_only	110.765	110.879	111.078	94.952	112.078	1.18x
programme_free	7.545	45.646	12.194	18.881	22.518	6.05x

$n = 3$ (seeds 42/43/44 only), as printed in paper/P2_FIGURES.md F3: programme_only 1.003x, programme_free 6.05x.

SUPERSEDED. Seeds 45 and 46 landed afterwards in `~/e0_run/d1_seeds4546/` and are included above; the stable arm's floor is 1.18x, not 1.003x.

F3. Statistic R3 - the in-run tripwire statistic, NOT R1 - and the block is a TRAINING batch at global step 200, not a held-out set. The states behind these numbers were never saved, so they are [NOT RECOMPUTABLE] under R1, and no number in this figure may share an axis with any R1 value in this paper. The key is ``train_rank_tripwire_observed`` with the ``train_`` prefix, because querying it without the prefix returns [] and an empty list reads as a confident negative. The two rows are not connected and share only the axis: they are two arms, not a series. The reading is that rank's reproducibility is worst exactly where the arm is interesting - which is the configuration a practitioner is looking at when they reach for rank. F3 and F4(a) are not in conflict: F4(a) holds the model fixed and varies the patient sample; F3 holds the step fixed and varies the seed.