




(a) Sampling noise with the trained model held FIXED
40 draws of 80% of patients per artifact



(c) Fixed seed, fixed 200-step horizon, identical inputs
three repeats of a controlled probe




(b) Readout determinism



deterministic to five significant figures
(0.58612 against 0.5861 recorded).
The export and readout path is not the source of the variance.

(d) Definitional insensitivity



Over all 68 recomputed artifact x block combinations:
absolute $\rightarrow$ relative tolerance change moved NO value (max relative difference 0.000e+00);
uncentred R1 / centred R1 has median 0.995 (min 0.116, max 1.000).
Source: NOTEBOOK_ENTRIES/effective_rank_canonicalised_..._20260804T0005Z.md, 2 and 4.

F4. In panel (a), D2 s44 is the ONE unresolvable pair - the between-arm gap is 1.4 sampling standard deviations - and section 4.6's D2 "2 of 3" depends on it. It is not a hit; it is a pair the metric cannot separate, which is why T1 marks that cell unresolvable in every row. Panel (c) constrains the TYPICAL spread and essentially not the tail: three repetitions cannot rule out the liveness gate's 25% divergence rate ($P(0 \text{ in } 3 \mid p = 0.25) = 0.42$; the exact upper 95% bound from 0 of 3 is $p \leq 0.63$), and the design was cut from ten repetitions to three because a ten-way launch exhausted GPU memory. Panel (c)'s statistic is R3 on a training probe and must not share an axis with (a)'s and (d)'s R1 on held-out artifacts; they are drawn on separate axes for that reason. Panel (a) and F3 are not in conflict: (a) holds the model fixed and varies the patient sample, F3 holds the step fixed and varies the seed. Panel (b) shares its data with F1(a). The faithful RankMe of panel (d) is computed on OUR OWN artifacts and is an internal comparison; no number in this paper is plotted against a published RankMe value, because RankMe's epsilon sits outside the division and its p_k do not sum to 1.