(a) The STATISTIC. Does higher rank pick the arm with more information?

	D2 s42	D2 s43	D2 s44	D1 s42	D1 s43	D1 s44
R1 canonical, order 1, centred	✓	✗	✓	✓	✓	✓
R2 order-2 Hill of the singular values	✓	✓	✗	✓	✓	✓
R3 order-2 Hill, rows L2-normalised	✓	✗	✓	✓	✓	✓
PR eigenvalue participation ratio	✓	✓	✗	✗	✓	✓
$PR_{rownorm}$ eigenvalue PR, rows normalised	✓	✓	✗	✓	✓	✓
INFORMATION verdict held-out top-CCA, 40 untrained targets	✓	✓	✓	✓	✓	✓

R1, R2 and R3 disagree on 2 of 6 pairs (D2 s43, D2 s44, boxed). The information verdict is correct in every column.
The lower block is the SAME data under a disambiguated name: the eigenvalue participation ratio, which an earlier version of this analysis computed under the R2 / R3 labels. PR is identical cell for cell to T1's participation-ratio row, so the two are NOT independent evidence.

(b) The BLOCK. Raw against confound-residualised, the two D2 seeds where R3 flips

		RAW block			RESIDUALISED block			
seed	stat	H	I	winner	H	I	winner	
43	R1	24.674	28.800	I	28.771	34.117	I	
43	R2	11.720	11.111	H	13.227	12.972	H	
43	R3	11.720	11.111	H	14.746	15.915	I	<-- FLIPS
44	R1	8.447	8.313	H	9.143	9.105	H	
44	R2	5.385	5.449	I	5.733	5.815	I	
44	R3	5.385	5.449	I	6.564	6.302	H	<-- FLIPS

The block matters for ONE statistic and not the others: R1 and R2 keep their ordering,
R3 reverses it on both seeds. On the RAW exported artifacts R2 and R3 are EQUAL to nine significant figures, because the model already L2-normalises z_biology - which is why the distinction went unnoticed. Statistic and block are named on every row; nothing here is an unqualified "effective rank".

(c) The VIEW. Three co-trained views of the same model

	D2 s42	D2 s43	D2 s44	D1 s42	D1 s43	D1 s44
rank winner, wsi_biology	H	I	H	P	P	P
rank winner, rna_biology	I	I	I	P	P	P
rank winner, full_biology	I	I	I	P	P	P
INFORMATION winner (identical on all three views)	H	H	H	P	P	P

Rank agrees with information on 11 of 18 (pair x view) comparisons; restricted to D2, 2 of 9.
Rank uses R1 (canonical effective rank, centred); information is the held-out top-CCA on the 40 untrained targets.

F5. These panels do not show that rank is wrong. They show that A RANK VERDICT IS NOT A WELL-DEFINED OBJECT until the statistic, the block and the view are stated - and that the three choices are worth more than the between-arm difference they are used to adjudicate. Panel (a) is drawn from ws_p2/out/P2_RANK_VARIANTS.json, not from any printed table, and its eigenvalue rows are labelled PR and PR_rownorm, never R2/R3: a fourth statistic lived under those labels inside the analysis code of the section that argues the name is unreliable, and was corrected at commit a11549a. PR is identical cell for cell to T1's participation-ratio row and is not independent evidence. In panel (c) the rna_biology -> RNA-derived-target comparison is partly circular and its absolute CCA (0.79-0.85) must not be read as a clean image-to-molecular channel; the RANK measurements on that view are unaffected by the circularity, but the D2 count inherits the caveat and is not quoted as a rate.